
1. Artificial Neurons

Structure and Mathematical Representation

An artificial neuron is inspired by biological neurons and consists of the following components:

- **Inputs (x_1, x_2, \dots, x_n):** Features or data points fed into the neuron.
- **Weights (w_1, w_2, \dots, w_n):** Each input has an associated weight that determines its importance.
- **Bias (b):** A constant added to shift the activation.
- **Summation Function:** Computes the weighted sum of inputs: $z = \sum (w_i x_i) + b$
- **Activation Function ($f(z)$):** Introduces non-linearity to decide whether the neuron should be "activated" or not.

Mathematical Representation

$$y = f\left(\sum_{i=1}^n w_i x_i + b\right)$$

Example

Let's implement a simple artificial neuron in Python:

```
import numpy as np
```

```
# Define inputs, weights, and bias
```

```
inputs = np.array([0.5, 0.3, 0.2])
```

```
weights = np.array([0.4, 0.7, 0.2])
```

```
bias = 0.1
```

```
# Compute the weighted sum
```

```
z = np.dot(inputs, weights) + bias
```

```
print(f"Summation (z): {z}")
```

```
# Apply activation function (Sigmoid)
```

```
activation = 1 / (1 + np.exp(-z))
```

```
print(f"Activated Output: {activation}")
```

2. Activation Functions

Activation functions introduce non-linearity into neural networks.

Types of Activation Functions

1. Linear Activation

$$f(z) = z$$

- Used in regression tasks.
- No non-linearity, so it's rarely used in deep networks.

2. Sigmoid Activation

$$f(z) = \frac{1}{1 + e^{-z}}$$

- Outputs values in the range (0,1).
- Suitable for binary classification.
- Problem: **Vanishing gradient** in deep networks.

```
import matplotlib.pyplot as plt
```

```
def sigmoid(z):
```

```
    return 1 / (1 + np.exp(-z))
```

```
z = np.linspace(-10, 10, 100)
```

```
plt.plot(z, sigmoid(z))
```

```
plt.title("Sigmoid Activation Function")
```

```
plt.grid()
```

```
plt.show()
```

3. Tanh Activation

$$f(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

- Output range (-1,1), centered around 0.
- Better than Sigmoid but still suffers from vanishing gradients.

```
def tanh(z):
```

```
    return np.tanh(z)
```

```
plt.plot(z, tanh(z))
plt.title("Tanh Activation Function")
plt.grid()
plt.show()
```

4. ReLU (Rectified Linear Unit)

$f(z) = \max(0, z)$

- Solves vanishing gradient for positive values.
- Computationally efficient.
- **Problem:** Neurons can "die" when $z < 0$ (Dying ReLU).

```
def relu(z):
    return np.maximum(0, z)
```

```
plt.plot(z, relu(z))
plt.title("ReLU Activation Function")
plt.grid()
plt.show()
```

5. Leaky ReLU

$f(z) = \max(0.01z, z)$

- Helps with the **Dying ReLU problem** by allowing small gradients when $z < 0$.

```
def leaky_relu(z, alpha=0.01):
    return np.where(z > 0, z, alpha * z)
```

```
plt.plot(z, leaky_relu(z))
plt.title("Leaky ReLU Activation Function")
plt.grid()
plt.show()
```

6. Softmax Activation

$f(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$

- Used in multi-class classification.

```
def softmax(z):  
    exp_z = np.exp(z - np.max(z))  
    return exp_z / exp_z.sum()  
  
z = np.array([1.0, 2.0, 3.0])  
print("Softmax output:", softmax(z))
```

3. ANN Architecture

Layers of an ANN

- **Input Layer:** Takes input features.
- **Hidden Layers:** Perform computations using activation functions.
- **Output Layer:** Produces final predictions.

Weight Initialization

- **Random Initialization** (Standard approach).
- **Xavier Initialization** (for Sigmoid/Tanh).
- **He Initialization** (for ReLU).

```
import tensorflow as tf  
  
# Xavier Initialization  
initializer = tf.keras.initializers.GlorotUniform()  
weights = initializer(shape=(3, 3))  
print(weights.numpy())
```

4. Forward and Backward Propagation

Forward Propagation

1. Compute weighted sum: $z = W \cdot X + b$
2. Apply activation function.

Backward Propagation (Learning Process)

- Uses **Gradient Descent** and **Chain Rule** to update weights.

Chain Rule for Gradients

If $y=f(g(x))$, then

$$\frac{dy}{dx} = f'(g(x)) \cdot g'(x) \quad \frac{dy}{dx} = f'(g(x)) \cdot g'(x)$$

Example

```
import torch

# Define loss function
loss_fn = torch.nn.MSELoss()

# Dummy data
y_pred = torch.tensor([0.4, 0.7, 0.2], requires_grad=True)
y_true = torch.tensor([0.5, 0.6, 0.3])

# Compute loss
loss = loss_fn(y_pred, y_true)
loss.backward() # Backpropagation

# Print gradients
print(y_pred.grad)
```

5. Training Neural Networks (Hands-on Implementation)

Let's implement a simple neural network using **TensorFlow (Keras)**.

Dataset: MNIST Digit Classification

```
import tensorflow as tf

from tensorflow.keras import layers, models

# Load dataset
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
```

```
# Normalize
```

```
x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
# Define model
```

```
model = models.Sequential([  
    layers.Flatten(input_shape=(28, 28)),  
    layers.Dense(128, activation='relu'),  
    layers.Dense(10, activation='softmax')  
])
```

```
# Compile model
```

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
# Train model
```

```
model.fit(x_train, y_train, epochs=5, validation_data=(x_test, y_test))
```

Using PyTorch

```
import torch
```

```
import torch.nn as nn
```

```
import torch.optim as optim
```

```
from torchvision import datasets, transforms
```

```
# Load data
```

```
train_loader = torch.utils.data.DataLoader(  
    datasets.MNIST('.', train=True, transform=transforms.ToTensor()),  
    batch_size=32, shuffle=True)
```

Define Model

```
class NeuralNet(nn.Module):
```

```
    def __init__(self):
```

```
        super().__init__()
```

```
        self.fc1 = nn.Linear(28*28, 128)
```

```
        self.fc2 = nn.Linear(128, 10)
```

```
    def forward(self, x):
```

```
        x = torch.relu(self.fc1(x.view(-1, 28*28)))
```

```
        return torch.softmax(self.fc2(x), dim=1)
```

Training

```
model = NeuralNet()
```

```
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
criterion = nn.CrossEntropyLoss()
```

```
for epoch in range(5):
```

```
    for images, labels in train_loader:
```

```
        optimizer.zero_grad()
```

```
        output = model(images)
```

```
        loss = criterion(output, labels)
```

```
        loss.backward()
```

```
        optimizer.step()
```
