

## Regularization Techniques in Deep Learning

When training deep learning models, **overfitting** is a common problem where the model performs well on training data but poorly on unseen test data. Regularization techniques help prevent overfitting by adding constraints to the model. Three widely used regularization techniques are **L1/L2 regularization, Dropout, and Early Stopping**.

---

### L1 & L2 Regularization (Weight Decay)

These are techniques that modify the **loss function** by adding a penalty term to discourage large weights, leading to simpler models that generalize better.

#### L1 Regularization (Lasso)

- Adds **absolute values** of weights to the loss function:  $LL1 = \sum |w_i|$   $L_{\text{L1}} = \sum |w_i|$
- Encourages **sparsity** (some weights become exactly 0), which can be useful for feature selection.
- Often used when feature selection is important.

#### L2 Regularization (Ridge)

- Adds **squared values** of weights to the loss function:  $LL2 = \sum w_i^2$   $L_{\text{L2}} = \sum w_i^2$
- Penalizes large weights, making them smaller but **not exactly zero**.
- Helps distribute weight values evenly, improving model stability.

#### Elastic Net (L1 + L2)

- Combines both L1 and L2 regularization:  $LElasticNet = \alpha LL1 + \beta LL2$   $L_{\text{ElasticNet}} = \alpha L_{\text{L1}} + \beta L_{\text{L2}}$
- Provides both sparsity and small weight values.

### Implementation in Python (TensorFlow/Keras)

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Dense
```

```
from tensorflow.keras.regularizers import l1, l2
```

```
model = Sequential([
```

```
    Dense(64, activation='relu', kernel_regularizer=l1(0.01)), # L1 Regularization
```

```
    Dense(64, activation='relu', kernel_regularizer=l2(0.01)), # L2 Regularization
```

```
    Dense(1, activation='sigmoid')
```

])

---

## 2 Dropout

Dropout is a **randomized** regularization technique where neurons are randomly dropped (set to zero) during training, preventing the network from relying too much on specific neurons.

### How It Works

- During each training step, a fraction of neurons is randomly set to zero.
- This forces the model to learn **multiple independent representations**, reducing overfitting.
- At test time, all neurons are used, but their outputs are scaled.

### Implementation in Python

```
from tensorflow.keras.layers import Dropout
```

```
model = Sequential([  
    Dense(64, activation='relu'),  
    Dropout(0.5), # 50% neurons dropped  
    Dense(64, activation='relu'),  
    Dropout(0.3), # 30% neurons dropped  
    Dense(1, activation='sigmoid')  
])
```

- Dropout rate (e.g., 0.5) is a **hyperparameter** that can be tuned.

---

## 3 Early Stopping

Early stopping **monitors validation performance** and stops training when the model starts overfitting.

### How It Works

- The model is trained for multiple epochs.
- A **monitoring metric** (like validation loss) is observed.
- If the validation loss stops improving for a set number of epochs (patience), training stops.

### Implementation in Python

```
from tensorflow.keras.callbacks import EarlyStopping
```

```
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
```

```
model.fit(X_train, y_train, epochs=100, validation_data=(X_val, y_val), callbacks=[early_stopping])
```

- `monitor='val_loss'`: Watches validation loss.
- `patience=5`: Waits for 5 epochs before stopping.
- `restore_best_weights=True`: Ensures the model returns to the best weights.

---

### Comparison of Techniques

Technique	Purpose	Effect on Weights
<b>L1 Regularization</b>	Reduces overfitting, feature selection	Some weights become <b>0</b> (sparse)
<b>L2 Regularization</b>	Reduces overfitting, stabilizes training	Shrinks all weights but keeps them nonzero
<b>Dropout</b>	Prevents co-adaptation of neurons	Randomly deactivates neurons during training
<b>Early Stopping</b>	Prevents overtraining	Stops training when validation performance deteriorates

---

### Which Regularization to Use?

- Use **L1 Regularization** if you need feature selection.
- Use **L2 Regularization** if you want to penalize large weights.
- Use **Dropout** for deep networks to prevent neuron dependence.
- Use **Early Stopping** when training for many epochs to avoid overfitting.