

## 1. Computer Vision Tasks

Computer vision is a field of AI that enables machines to interpret and process visual data. The key tasks in computer vision include:

### 1.1 Image Classification

- **Definition:** Assigning a label to an image from a predefined set of categories.
- **Example:** Classifying images into "cat" or "dog."

### Mathematical Explanation

Given an input image  $X$ , a classification model outputs a probability distribution over  $C$  classes:

$$P(y|X) = f(X, \theta) \quad P(y | X) = f(X, \theta)$$

where  $f$  is a deep neural network with parameters  $\theta$ , and  $y$  is the predicted class.

The model is trained using **Cross-Entropy Loss**:

$$L = -\sum_{i=1}^C y_i \log(\hat{y}_i) \quad L = -\sum_{i=1}^C y_i \log(\hat{y}_i)$$

where  $y_i$  is the true label and  $\hat{y}_i$  is the predicted probability.

### Python Code for Image Classification using CNN

```
import tensorflow as tf

from tensorflow import keras

from tensorflow.keras import layers

# Load dataset (CIFAR-10)

(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()

x_train, x_test = x_train / 255.0, x_test / 255.0 # Normalize images

# Define CNN model

model = keras.Sequential([

    layers.Conv2D(32, (3,3), activation='relu', input_shape=(32,32,3)),

    layers.MaxPooling2D((2,2)),

    layers.Conv2D(64, (3,3), activation='relu'),

    layers.MaxPooling2D((2,2)),

    layers.Flatten(),
```

```

layers.Dense(64, activation='relu'),
layers.Dense(10, activation='softmax') # 10 classes in CIFAR-10
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Train model
model.fit(x_train, y_train, epochs=10, validation_data=(x_test, y_test))

```

---

## 1.2 Object Detection

- **Definition:** Identifying and locating objects in an image by drawing bounding boxes around them.
- **Example:** Detecting multiple objects like "person" and "car" in an image.

### Mathematical Explanation

Object detection involves:

- **Classification** (which object is present)
- **Localization** (where the object is)

A neural network predicts:

$(x, y, w, h, c_1, c_2, \dots, c_C)$   $(x, y, w, h, c_1, c_2, \dots, c_C)$

where:

- $(x, y)$   $(x, y)$  is the center of the bounding box.
- $(w, h)$   $(w, h)$  is the width and height.
- $c_i$  are the class probabilities.

The **Intersection over Union (IoU)** metric evaluates detection accuracy:

$$\text{IoU} = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$

### Python Code for Object Detection using YOLO

```
from ultralytics import YOLO
```

```
# Load YOLOv8 pre-trained model

model = YOLO("yolov8n.pt") # Download model automatically


# Run inference on an image

results = model("image.jpg")


# Display results

results.show()
```

---

### 1.3 Image Segmentation

- **Definition:** Assigning each pixel of an image to a specific class.
- **Example:** Segmenting "sky," "road," and "car" in a driving scene.

#### Mathematical Explanation

Segmentation models output a probability map  $P$  of shape  $(H, W, C)$ , where each pixel gets a class label:

$$\hat{y}_{i,j} = \arg\max_c P_{i,j,c} \quad \hat{y}_{i,j} = \arg\max_c P_{i,j,c}$$

Loss function: **Dice Coefficient Loss**

$$L = 1 - \frac{2 \sum (P \cdot Y)}{\sum P + \sum Y}$$

#### Python Code for Image Segmentation using U-Net

```
import tensorflow as tf

from tensorflow.keras.layers import Conv2D, MaxPooling2D, UpSampling2D, concatenate


# Define U-Net architecture

def unet_model():

    inputs = keras.Input(shape=(128, 128, 3))

    conv1 = Conv2D(64, (3, 3), activation="relu", padding="same")(inputs)

    pool1 = MaxPooling2D((2, 2))(conv1)
```

```
conv2 = Conv2D(128, (3, 3), activation="relu", padding="same")(pool1)
up1 = UpSampling2D((2, 2))(conv2)
```

```
outputs = Conv2D(1, (1, 1), activation="sigmoid")(up1)
return keras.Model(inputs, outputs)
```

```
model = unet_model()
model.compile(optimizer="adam", loss="binary_crossentropy", metrics=["accuracy"])
```

---

## 2. Image Data Handling

Before training, images must be **loaded, resized, and normalized**.

### 2.1 Loading and Resizing Images

```
import cv2
import numpy as np
```

```
# Load image
image = cv2.imread("image.jpg")
```

```
# Resize to 224x224
resized_image = cv2.resize(image, (224, 224))
```

### 2.2 Normalizing Images

```
# Convert to float and normalize to range [0,1]
normalized_image = resized_image / 255.0
```

---

## 3. Data Augmentation

Improves generalization by applying transformations like rotation, flipping, cropping, and color jittering.

### 3.1 Rotation

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
datagen = ImageDataGenerator(rotation_range=30)
```

### 3.2 Flipping

```
datagen = ImageDataGenerator(horizontal_flip=True, vertical_flip=True)
```

### 3.3 Cropping

```
datagen = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1)
```

### 3.4 Color Jittering

```
datagen = ImageDataGenerator(brightness_range=[0.5, 1.5])
```

### 3.5 Full Data Augmentation Example

```
datagen = ImageDataGenerator(  
    rotation_range=30,  
    width_shift_range=0.1,  
    height_shift_range=0.1,  
    horizontal_flip=True,  
    brightness_range=[0.5, 1.5]  
)
```

```
# Apply to an image
```

```
image = np.expand_dims(image, 0) # Expand dimensions for batch
```

```
augmented_image = datagen.flow(image, batch_size=1)
```

---

## Conclusion

- **Image Classification** assigns a single label to an image.
- **Object Detection** localizes multiple objects using bounding boxes.
- **Image Segmentation** assigns each pixel a class.
- **Image Data Handling** includes loading, resizing, and normalizing.
- **Data Augmentation** improves generalization.