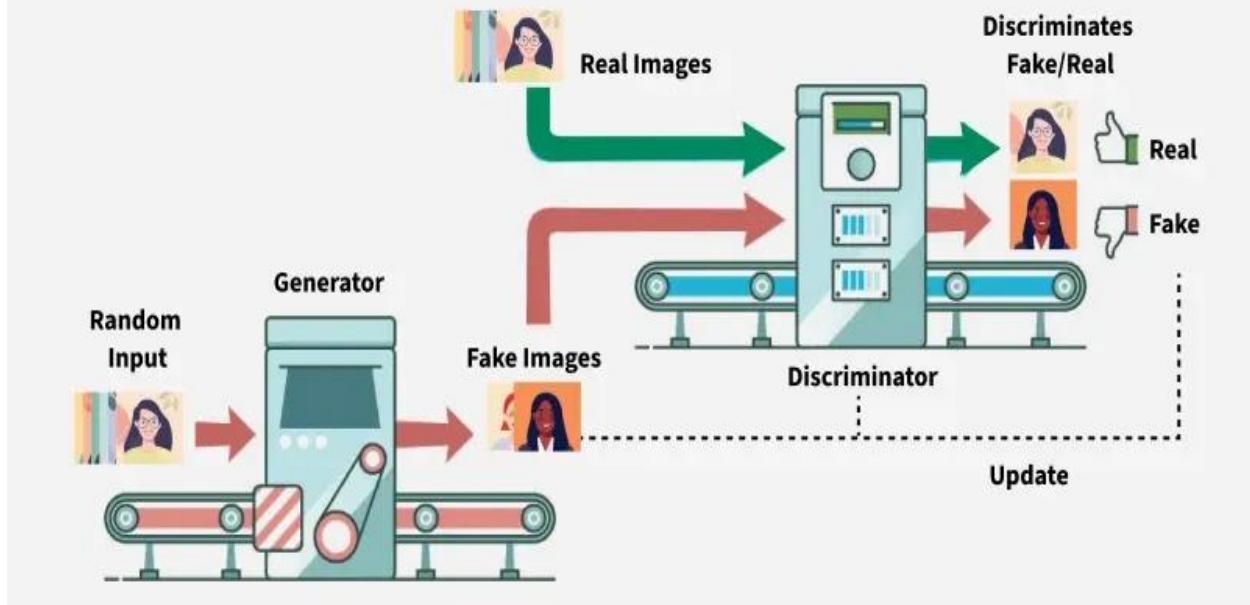


Generative Adversarial Network (GANs)



Generative Adversarial Networks (GANs) in Computer Vision

Learning Objectives

By the end of this lesson, you will:

1. Understand the architecture of GANs.
2. Learn how the Generator and Discriminator networks interact.
3. Explore the challenges in training GANs and best practices.
4. Implement a simple GAN for image generation in Python.

1. GAN Architecture

Generative Adversarial Networks (GANs) were introduced by Ian Goodfellow in 2014. GANs consist of two neural networks that compete with each other in a game-theoretic framework:

1.1 Generator (G)

- The Generator takes a random noise vector z as input.
- It learns to generate realistic-looking images.
- The goal is to **fool the Discriminator** into classifying the generated images as real.

Mathematically, the Generator $G(z)$ maps a random noise z (sampled from a distribution, such as a normal distribution) to the data space (images):

$$G(z; \theta_G) \rightarrow X_{\text{fake}} \quad G(z; \theta_G) \rightarrow X_{\text{fake}}$$

where θ_G are the parameters of the Generator.

1.2 Discriminator (D)

- The Discriminator takes an image as input and predicts whether it's real or fake.
- It is a binary classifier trained using cross-entropy loss.
- It learns to distinguish between real images from the dataset and fake images from the Generator.

Mathematically, the Discriminator $D(X)$ outputs a probability that the input X is real:

$$D(X; \theta_D) \rightarrow [0, 1] \quad D(X; \theta_D) \rightarrow [0, 1]$$

where θ_D are the parameters of the Discriminator.

2. Training GANs

2.1 Objective Function (Minimax Game)

The Generator and Discriminator are trained simultaneously using the following loss function:

$$\min_G \max_D V(D, G) = \mathbb{E}_{X \sim p_{\text{data}}} [\log D(X)] + \mathbb{E}_{Z \sim p_Z} [\log (1 - D(G(Z)))]$$

$$\min_G \max_D V(D, G) = \mathbb{E}_{X \sim p_{\text{data}}} [\log D(X)] + \mathbb{E}_{Z \sim p_Z} [\log (1 - D(G(Z)))]$$

- The **Discriminator** tries to maximize $\log D(X) + \log (1 - D(G(Z)))$ (it gets better at distinguishing real and fake images).
- The **Generator** tries to minimize $\log (1 - D(G(Z)))$ (it tries to fool the Discriminator).

2.2 Training Steps

1. Train the Discriminator:

- Feed real images from the dataset and train D to classify them as real.
- Feed generated images from G and train D to classify them as fake.

2. Train the Generator:

- Generate fake images using random noise.
- Compute the loss based on how well the Discriminator is fooled.
- Update G to produce more realistic images.

2.3 Challenges in Training

Training GANs is difficult due to:

- **Mode Collapse:** The Generator produces limited variations of images.
- **Vanishing Gradients:** The Discriminator may become too strong, making it hard for the Generator to learn.
- **Instability:** Training is highly sensitive to hyperparameters.

2.4 Best Practices

- Use **batch normalization** in both networks.
 - Use **Leaky ReLU** activation in the Discriminator.
 - Apply **label smoothing** to make training more stable.
 - Use **Adam optimizer** with proper hyperparameters.
-

3. Implementation: Building a Simple GAN in Python

We will use TensorFlow and Keras to implement a GAN for generating handwritten digits using the **MNIST dataset**.

Step 1: Import Libraries

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import numpy as np
import matplotlib.pyplot as plt
```

Step 2: Build the Generator

```
def build_generator():
    model = keras.Sequential([
        layers.Dense(128, activation="relu", input_shape=(100,)),
        layers.BatchNormalization(),
        layers.Dense(256, activation="relu"),
        layers.BatchNormalization(),
```

```
layers.Dense(512, activation="relu"),
layers.BatchNormalization(),
layers.Dense(28 * 28, activation="tanh"),
layers.Reshape((28, 28))
])
return model
```

- Input: 100-dimensional noise vector.
- Output: 28×28 grayscale image.
- Uses **Batch Normalization** for stable training.

Step 3: Build the Discriminator

```
def build_discriminator():
    model = keras.Sequential([
        layers.Flatten(input_shape=(28, 28)),
        layers.Dense(512, activation="leaky_relu"),
        layers.Dense(256, activation="leaky_relu"),
        layers.Dense(1, activation="sigmoid") # Binary classification
    ])
    return model
```

- Input: 28×28 grayscale image.
- Output: Probability (0 = fake, 1 = real).
- Uses **Leaky ReLU** for better gradient flow.

Step 4: Compile the GAN

```
def compile_gan(generator, discriminator):
    discriminator.compile(loss="binary_crossentropy", optimizer=keras.optimizers.Adam(0.0002, 0.5))

    gan = keras.Sequential([generator, discriminator])
    discriminator.trainable = False # Freeze Discriminator while training Generator
```

```
gan.compile(loss="binary_crossentropy", optimizer=keras.optimizers.Adam(0.0002, 0.5))
```

```
return gan
```

Step 5: Train the GAN

```
def train_gan(generator, discriminator, gan, epochs=10000, batch_size=128):
```

```
    (X_train, _) = keras.datasets.mnist.load_data()
```

```
    X_train = (X_train / 127.5) - 1 # Normalize images to [-1, 1]
```

```
    for epoch in range(epochs):
```

```
        # Train Discriminator
```

```
        real_images = X_train[np.random.randint(0, X_train.shape[0], batch_size)]
```

```
        noise = np.random.normal(0, 1, (batch_size, 100))
```

```
        fake_images = generator.predict(noise)
```

```
        X = np.concatenate([real_images, fake_images])
```

```
        y = np.concatenate([np.ones((batch_size, 1)), np.zeros((batch_size, 1))])
```

```
        d_loss = discriminator.train_on_batch(X, y)
```

```
        # Train Generator
```

```
        noise = np.random.normal(0, 1, (batch_size, 100))
```

```
        y_fake = np.ones((batch_size, 1)) # Trick Discriminator
```

```
        g_loss = gan.train_on_batch(noise, y_fake)
```

```
        # Print progress
```

```
        if epoch % 1000 == 0:
```

```
            print(f"Epoch {epoch} | D Loss: {d_loss:.4f} | G Loss: {g_loss:.4f}")
```

```
train_gan(build_generator(), build_discriminator(), compile_gan(build_generator(),
build_discriminator()))
```

Step 6: Generate Images

```
def generate_images(generator, n=10):
    noise = np.random.normal(0, 1, (n, 100))
    images = generator.predict(noise)
    images = (images + 1) / 2 # Rescale to [0,1]

    plt.figure(figsize=(10, 2))
    for i in range(n):
        plt.subplot(1, n, i + 1)
        plt.imshow(images[i], cmap="gray")
        plt.axis("off")
    plt.show()
```

```
generator = build_generator()
discriminator = build_discriminator()
gan = compile_gan(generator, discriminator)

train_gan(generator, discriminator, gan, epochs=10000, batch_size=128)
generate_images(generator)
```

Conclusion

- **GANs** consist of a Generator and a Discriminator in a competitive setting.
 - **Training is challenging**, but best practices like Batch Normalization and Leaky ReLU help.
 - **We implemented a simple GAN** to generate handwritten digits using TensorFlow/Keras.
-

Video resource: <https://youtu.be/FelMzID2BNg?si=rLX6ytS0hzEi43Ps>